


Rotman Research Node (RRN)


Jay / TDMDAL

What is RNN (1)



- A shared research server
 - Hardware: 32 Xeon CPUs; 800G Memory
 - Storage: Home directories; 50T project directory
 - Software: Python, R, Matlab, Stata, Julia, C, C++, Fortran, etc.
 - Dedicated to Rotman researchers
 - Hosted at Centre for Advanced Computing (CAC) at Queens U
- A gateway to a pool of computing resources
 - Zoom out (next slides)



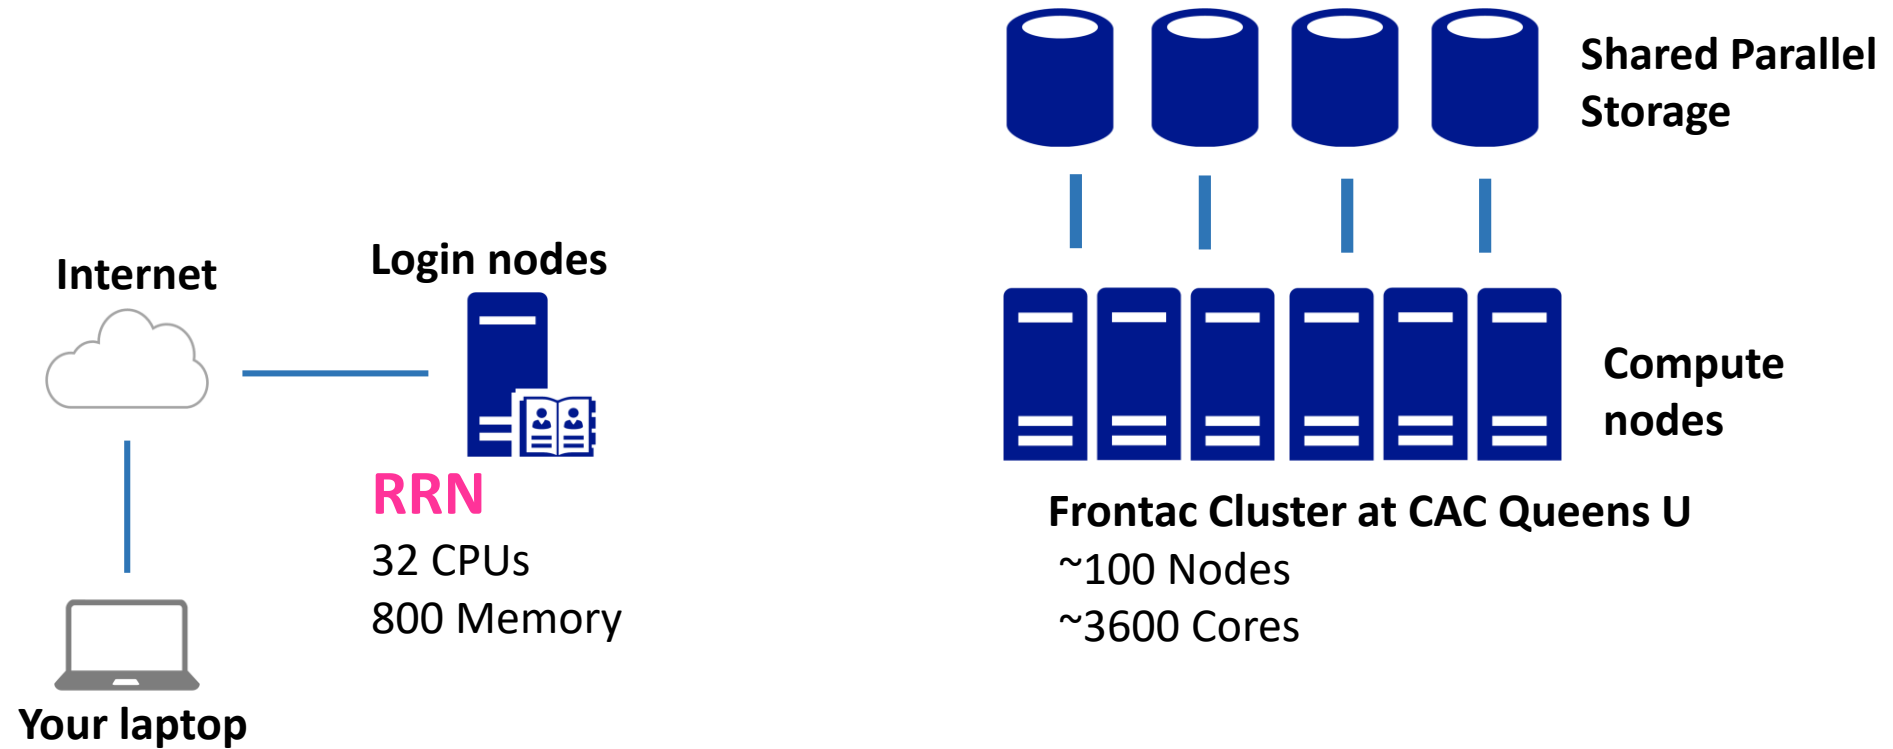
What's RRN (Zoom Out)



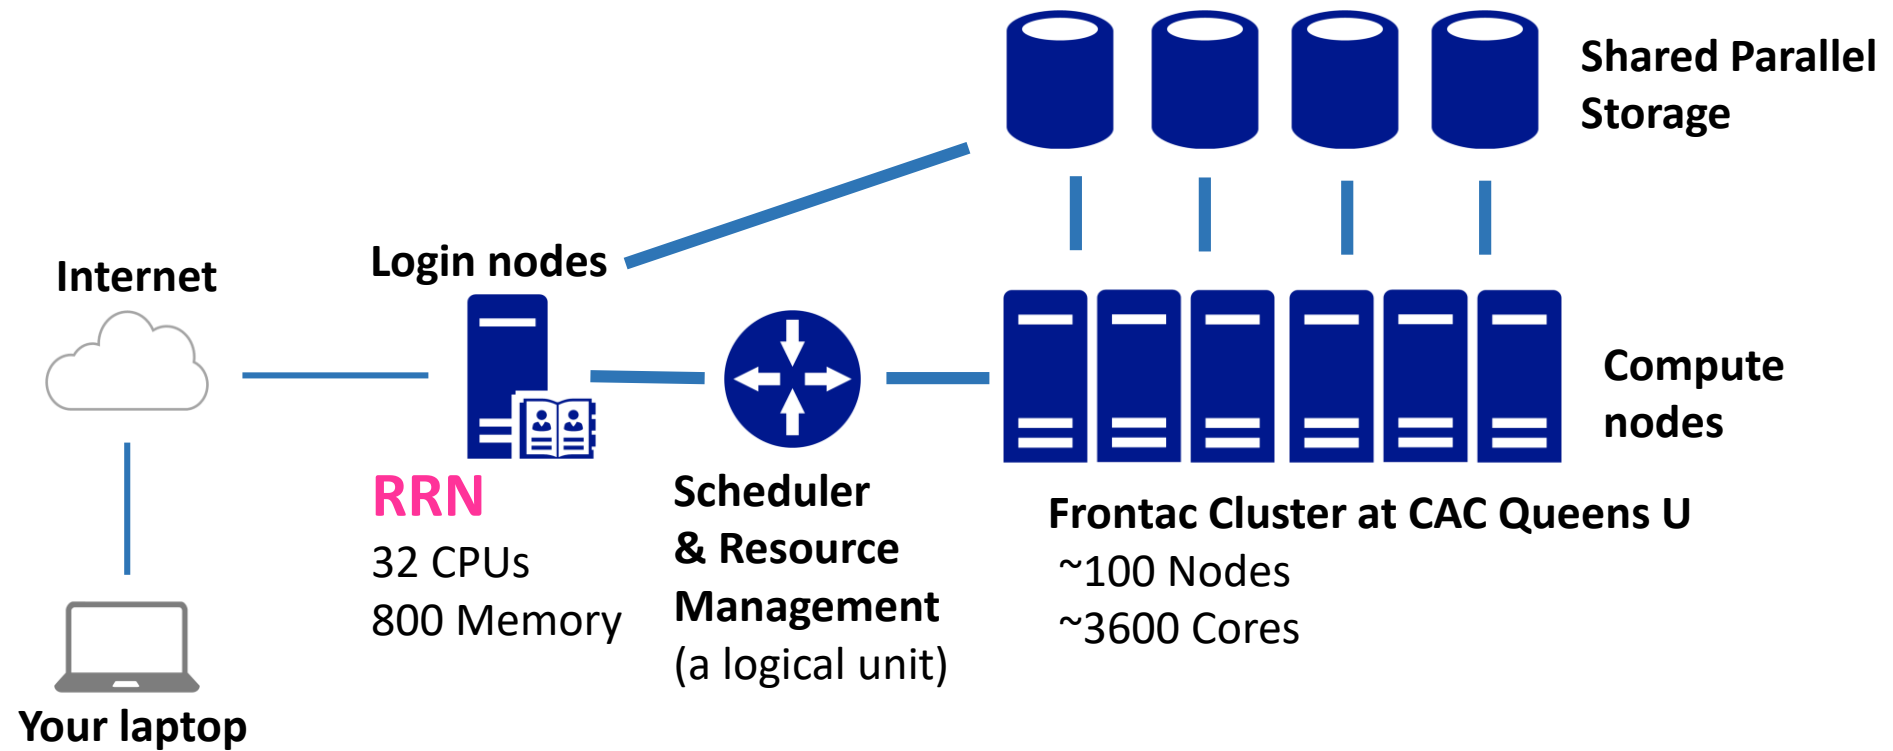
What's RRN (Zoom Out)



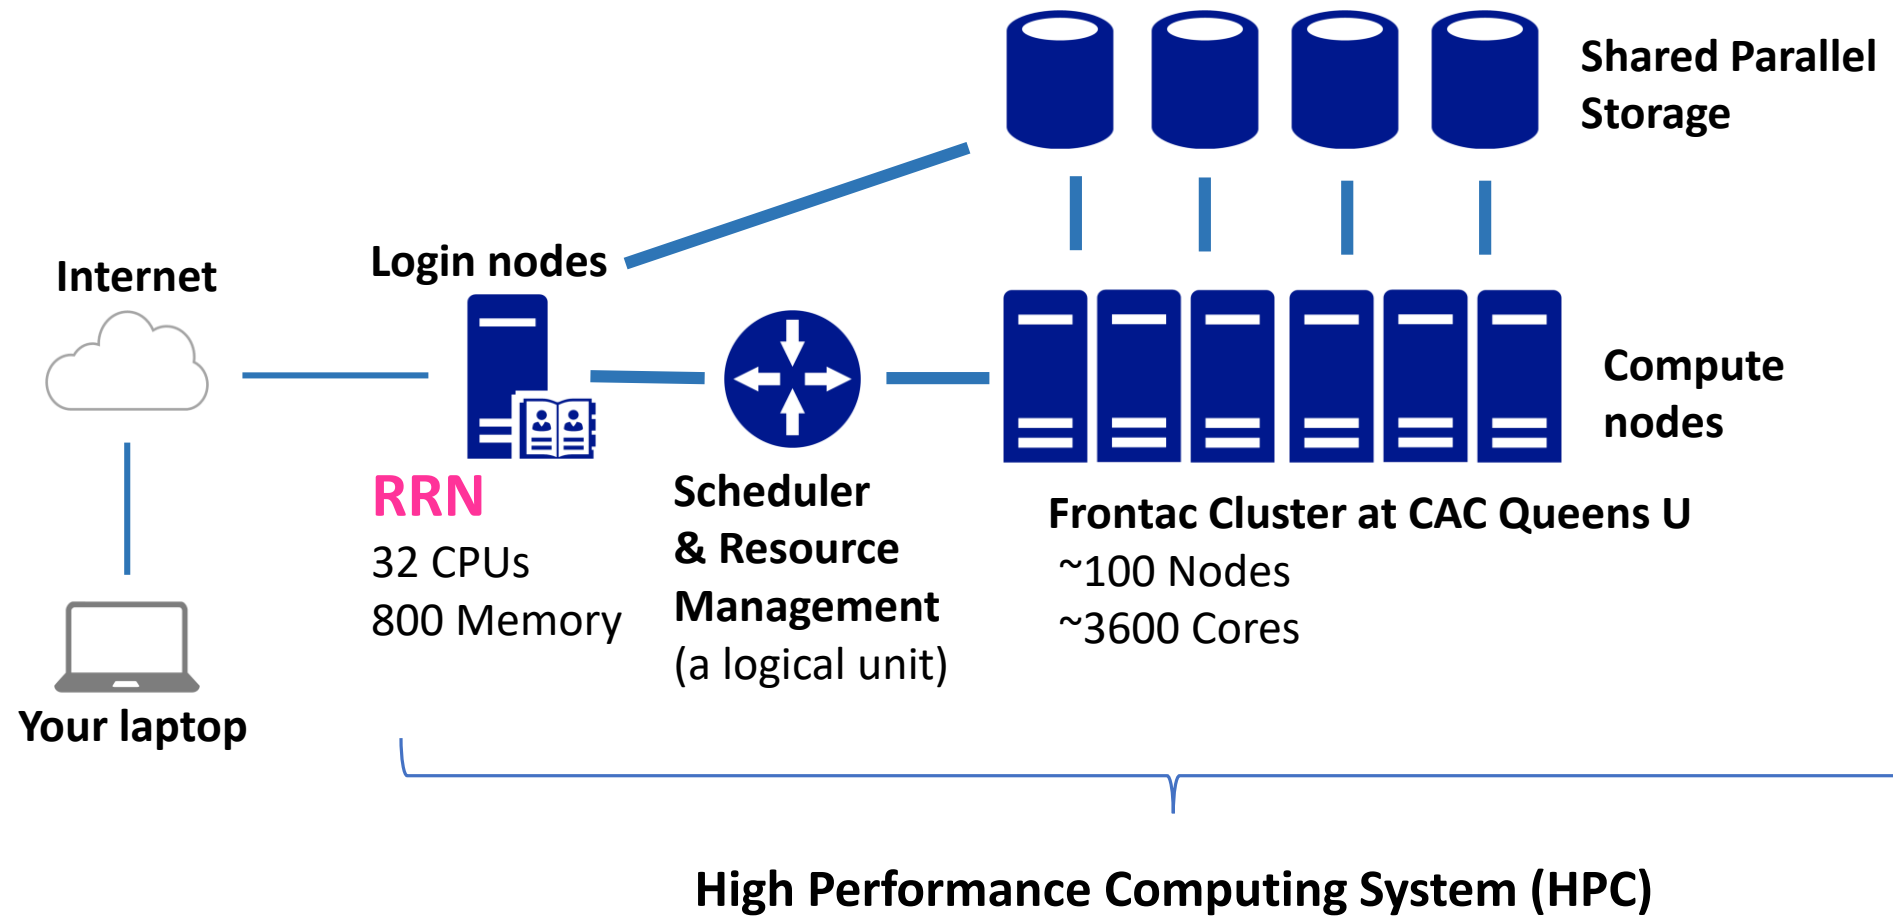
What's RRN (Zoom Out)



What's RRN (Zoom Out)



What's RRN (Zoom Out)



Why using RRN

- My code takes too long to run
 - it need more CPUs (i.e. computing intensive work)
 - it need more memory (i.e. memory intensive work)
- I want to use XYZ (ex. Matlab), but the license cost is too high
- One big misunderstanding
 - my (unmodified) code will run faster on the server (**NO** in most cases)

Questions to ask before considering RRN

- I need more CPUs
 - have I optimized my code (vectorization; better algorithm)?
 - have I tried parallel computing on my desktop?
 - modern desktop has 2-4 CPUs (4-8 with hyperthreading)
 - no license cost if you use open source language (R, Python, Julia, etc.)
- I need more memory
 - do I really need to load all those data
 - have I optimized my code (delete big variables/objects after use; better algorithm)

How to Use RRN – User Manuals

- Our RRN [User Manual](#)
 - Focus on RRN
 - Good for getting started
- CAC [User manual](#)
 - Including HPC usage (Compute Nodes)
 - For users with highly compute-intensive jobs

How to Use RRN – Account & Logon

- Email tdmdal@rotman.utoronto.ca for an account
- SSH client
 - Windows (ex. [Mobaxterm](#))
 - Mac (terminal + [Xquartz](#))

```
ssh -X yourUserName@rrlogin.cac.queensu.ca
```

How to Use RRN – Transfer Files

- a SFTP client (transfer files between local PC and RRN)
 - Windows (ex. [WinSCP](#), [Cyberduck](#), [FileZilla](#))
 - Mac ([Cyberduck](#), [FileZilla](#))
- Other methods available too
 - scp, rsync, etc.
 - globus
 - ...

<https://tdmdal.github.io/computing-research/getting-started.html#transferring-files>

<https://cac.queensu.ca/wiki/index.php/UploadingFiles:Frontenac>

When to Use What

	Interactive Mode	Batch Mode
RRN (1 node)		
Compute Nodes (Many nodes)		

When to Use What

	Interactive Mode	Batch Mode
RRN (1 node) <ul style="list-style-type: none">• CPU: 32 cores• Memory: 800G; quite large• Shared (among Rotman researchers)• Easy to use: Code runs as soon as you ask it to run		
Compute Nodes (Many nodes)		

When to Use What

	Interactive Mode	Batch Mode
RRN (1 node) <ul style="list-style-type: none">• CPU: 32 cores• Memory: 800G; quite large• Shared (among Rotman researchers)• Easy to use: Code runs as soon as you ask it to run		
Compute Nodes (Many nodes) <ul style="list-style-type: none">• CPU: 24- up to 128-core node• Memory: mostly 256G nodes; a few large ones (512G, 1T & 2T)• Exclusive use once allocated to you• Slightly harder to use<ul style="list-style-type: none">• Need to write a script to talk to the scheduler/resource manager• Need to wait in a queue to compete for resources		

When to Use What

	Interactive Mode	Batch Mode
RRN (1 node) <ul style="list-style-type: none">• CPU: 32 cores• Memory: 800G; quite large• Shared (among Rotman researchers)• Easy to use: Code runs as soon as you ask it to run	Debug code Run small jobs <ul style="list-style-type: none">• < 12 cores• < 0.5 hrs.	
Compute Nodes (Many nodes) <ul style="list-style-type: none">• CPU: 24- up to 128-core node• Memory: mostly 256G nodes; a few large ones (512G, 1T & 2T)• Exclusive use once allocated to you• Slightly harder to use<ul style="list-style-type: none">• Need to write a script to talk to the scheduler/resource manager• Need to wait in a queue to compete for resources		

When to Use What

	Interactive Mode	Batch Mode
RRN (1 node) <ul style="list-style-type: none">• CPU: 32 cores• Memory: 800G; quite large• Shared (among Rotman researchers)• Easy to use: Code runs as soon as you ask it to run	Debug code Run small jobs <ul style="list-style-type: none">• < 12 cores• < 0.5 hrs.	Run intermediate compute & memory intensive jobs <ul style="list-style-type: none">• ~12 cores• 1-2 hrs. or ...
Compute Nodes (Many nodes) <ul style="list-style-type: none">• CPU: 24- up to 128-core node• Memory: mostly 256G nodes; a few large ones (512G, 1T & 2T)• Exclusive use once allocated to you• Slightly harder to use<ul style="list-style-type: none">• Need to write a script to talk to the scheduler/resource manager• Need to wait in a queue to compete for resources		

When to Use What

	Interactive Mode	Batch Mode
RRN (1 node) <ul style="list-style-type: none">• CPU: 32 cores• Memory: 800G; quite large• Shared (among Rotman researchers)• Easy to use: Code runs as soon as you ask it to run	Debug code Run small jobs <ul style="list-style-type: none">• < 12 cores• < 0.5 hrs.	Run intermediate compute & memory intensive jobs <ul style="list-style-type: none">• ~12 cores• 1-2 hrs. or ...
Compute Nodes (Many nodes) <ul style="list-style-type: none">• CPU: 24- up to 128-core node• Memory: mostly 256G nodes; a few large ones (512G, 1T & 2T)• Exclusive use once allocated to you• Slightly harder to use<ul style="list-style-type: none">• Need to write a script to talk to the scheduler/resource manager• Need to wait in a queue to compete for resources	Get a interactive node to: Debug code Run small jobs <ul style="list-style-type: none">• CPUs: as many as you request on a node• < 1 hrs.	

When to Use What

	Interactive Mode	Batch Mode
RRN (1 node) <ul style="list-style-type: none">• CPU: 32 cores• Memory: 800G; quite large• Shared (among Rotman researchers)• Easy to use: Code runs as soon as you ask it to run	Debug code Run small jobs <ul style="list-style-type: none">• < 12 cores• < 0.5 hrs.	Run intermediate compute & memory intensive jobs <ul style="list-style-type: none">• ~12 cores• 1-2 hrs. or ...
Compute Nodes (Many nodes) <ul style="list-style-type: none">• CPU: 24- up to 128-core node• Memory: mostly 256G nodes; a few large ones (512G, 1T & 2T)• Exclusive use once allocated to you• Slightly harder to use<ul style="list-style-type: none">• Need to write a script to talk to the scheduler/resource manager• Need to wait in a queue to compete for resources	Get a interactive node to: Debug code Run small jobs <ul style="list-style-type: none">• CPUs: as many as you request on a node• < 1 hrs.	Run highly compute & memory intensive jobs

Using Software – Module System

- A software Environment Module system

```
module avail
```

```
module load <module_name>
```

```
module list
```

Ex.1 Matlab - RNN Interactive Mode

- Load matlab module

```
module load matlab/R2018b
```

- Run matlab with GUI

```
matlab
```

- Run matlab without GUI

```
matlab -nodesktop -nosplash -nodisplay
```

Ex.1 Matlab - RNN Batch Mode

- Load matlab module

```
module load matlab/R2018b
```

- Run in foreground (i.e. prompt taken; need to waiting for result)

```
matlab -nodesktop -nosplash -nodisplay <matlab_test.m  
&>matlab_test.log
```

- Run in background and no hang up after logout

```
nohup matlab -nodesktop -nosplash -nodisplay  
<matlab_test.m &>matlab_test.log &
```

Ex.1 Matlab – Interactive Compute Node

- Allocate an interactive node

```
salloc -c 4 --mem=8g
```

- Load matlab module

```
module load matlab/R2018b
```

- Run matlab in interactive mode with or without GUI
- Run matlab in batch mode. **However, don't exit the interactive node!**

Ex.1 Matlab – Batch Compute Node (1)

- Write a job submission script (an example, job.sh)

use bash shell to
execute this script

```
#!/bin/bash
```

Slurm
command

```
#SBATCH --job-name=MATLAB_test  
#SBATCH --partition=standard  
#SBATCH --mail-type=ALL  
#SBATCH --mail-user=jay.cao@rotman.utoronto.ca  
#SBATCH --output=STD.out  
#SBATCH --error=STD.err  
#SBATCH -c 4  
#SBATCH --time=30:00  
#SBATCH --mem=5000
```

load matlab &
run my code

```
module load matlab/R2018b  
matlab -nodesktop -nosplash -nodisplay <matlab_test.m
```


Ex.1 Matlab – Batch Compute Node (2)

```
#!/bin/bash

#SBATCH --job-name=MATLAB_test      # set job name
#SBATCH --partition=standard        # set job partition (group of nodes)
#SBATCH --mail-type=ALL              # email me when job start, stop, etc.
#SBATCH --mail-user=jay.cao@rotman.utoronto.ca
#SBATCH --output=STD.out             # save standard output to STD.out
#SBATCH --error=STD.err              # save std. error output to STD.out
#SBATCH -c 4                         # ask for 4 CPUs
#SBATCH --time=30:00                 # set wall time to be 30mins
#SBATCH --mem=5000                   # ask for 5G memory

module load matlab/R2018b
matlab -nodesktop -nosplash -nodisplay <matlab_test.m
```

Ex.1 Matlab – Batch Compute Node (2)

```
#!/bin/bash

#SBATCH --job-name=MATLAB_test      # set job name
#SBATCH --partition=standard        # set job partition (group of nodes)
#SBATCH --mail-type=ALL             # email me when job start, stop, etc.
#SBATCH --mail-user=jay.cao@rotman.utoronto.ca # set my email address
#SBATCH --output=STD.out            # save standard output to STD.out
#SBATCH --error=STD.err             # save std. error output to STD.out
#SBATCH -c 4                        # ask for 4 CPUs
#SBATCH --time=30:00                # set wall time to be 30mins
#SBATCH --mem=5000                  # ask for 5G memory

module load matlab/R2018b
matlab -nodesktop -nosplash -nodisplay <matlab_test.m
```

Ex.1 Matlab – Batch Compute Node (2)

```
#!/bin/bash

#SBATCH --job-name=MATLAB_test      # set job name
#SBATCH --partition=standard        # set job partition (group of nodes)
#SBATCH --mail-type=ALL              # email me when job start, stop, etc.
#SBATCH --mail-user=jay.cao@rotman.utoronto.ca # set my email address
#SBATCH --output=STD.out             # save standard output to STD.out
#SBATCH --error=STD.err              # save std. error output to STD.out
#SBATCH -c 4                         # ask for 4 CPUs
#SBATCH --time=30:00                 # set wall time to be 30mins
#SBATCH --mem=5000                   # ask for 5G memory

module load matlab/R2018b
matlab -nodesktop -nosplash -nodisplay <matlab_test.m
```

Ex.1 Matlab – Batch Compute Node (2)

```
#!/bin/bash

#SBATCH --job-name=MATLAB_test      # set job name
#SBATCH --partition=standard        # set job partition (group of nodes)
#SBATCH --mail-type=ALL              # email me when job start, stop, etc.
#SBATCH --mail-user=jay.cao@rotman.utoronto.ca # set my email address
#SBATCH --output=STD.out            # save standard output to STD.out
#SBATCH --error=STD.err             # save std. error output to STD.out
#SBATCH -c 4                        # ask for 4 CPUs
#SBATCH --time=30:00                # set wall time to be 30mins
#SBATCH --mem=5000                  # ask for 5G memory

module load matlab/R2018b
matlab -nodesktop -nosplash -nodisplay <matlab_test.m
```

Ex.1 Matlab – Batch Compute Node (3)

- run job script to submit your code to a compute node
 - currently CAC doesn't enable inter-node jobs

```
sbatch job.sh
```

- Show status of jobs

```
queue --job <job_id>
```

- Cancel jobs

```
scancel <job_id>
```